

CS118 Project 1 – Web Server Implementation Using BSD Sockets

Anthony Lai 004445644

Elton Leong 204457607

High Level Description

The server design is based on the ClientServer_Example we were given to demonstrate how to use BSD sockets.

The server runs forever and forks itself whenever there is an incoming request. The child process then handles the incoming request while the parent remains to listen on the given port. There is a function `void respondToClient(int sockfd)` that processes and handles each request made to the server. It is called once by each child process. Most of the modifications to the code consist of changes to the function that handles the request.

After reading 511 bytes from the TCP socket, the HTTP request is parsed. The first line consisting of the method, URI, and version numbers is decoded, and this URI is saved for later, to lookup the file to be served. All of the HTTP request headers are parsed, stored in a hash table, and then ignored, as this is a rudimentary HTTP server.

For file lookup in the same directory as the webserver executable, a `.` is prepended to the request URI (e.g. `“.” + “/test.jpg” = “./test.jpg”`). If this file cannot be found, then a 404 Not Found is the response given by the server. Otherwise, a 200 OK is sent, along with the Content-Type and Content-Length headers, followed by the requested file.

Using a hash table, a hard coded mapping of file extensions to MIME types is made. For each request made to the server, the lower-case form of the extension is looked up in the hash table. If it cannot be found, the default MIME type `“application/octet-stream”` is given, which instead causes most browsers to download the given file.

Following the writing of all of the data to the TCP socket, the child process closes the socket fd and the child process exits successfully.

Difficulties

The only difficulty encountered was determining the minimum valid 404 response. A response consisting just of `“HTTP/1.1 404 Not Found\r\n”` caused Chrome to report that the connection terminated. Adding a basic `“Content-Length: 0”` header however satisfied Chrome.

Manual

To compile the code, type `make`

Run the server using `./webserver $portnum`, where `portnum` is the desired port, e.g.
`./webserver 42069`

To exit the server, hit `Ctrl+C`.

Sample Output

HTTP Request 1 (obtained from running the server and making a request using a web browser), a successful request of an HTML page

```
elton@TUESDAYSPECIAL:/mnt/c/Users/elton/Desktop/vm-shared/cs118/project1_204457607$ ./webserver 42069
```

```
GET /test.html HTTP/1.1
```

```
Accept: text/html, application/xhtml+xml, image/jxr, */*
```

```
Accept-Language: en-US,en;q=0.5
```

```
User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko)
Chrome/51.0.2704.79 Safari/537.36 Edge/14.14393
```

```
Accept-Encoding: gzip, deflate
```

```
DNT: 1
```

```
Host: localhost:42069
```

```
Connection: Keep-Alive
```

HTTP Response 1 (obtained from using Fiddler, a traffic snooper)

```
HTTP/1.1 200 OK
```

```
Content-Type: text/html
```

```
Content-Length: 25
```

```
<html>
```

```
    Hello
```

```
</html>
```

HTTP Request 2, a request of a file that doesn't exist

```
elton@TUESDAYSPECIAL:/mnt/c/Users/elton/Desktop/vm-shared/cs118/project1_204457607$ ./webserver 42069
```

```
GET /test HTTP/1.1
```

```
Accept: text/html, application/xhtml+xml, image/jxr, */*
```

```
Accept-Language: en-US,en;q=0.5
```

```
User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko)
Chrome/51.0.2704.79 Safari/537.36 Edge/14.14393
```

```
Accept-Encoding: gzip, deflate
```

```
DNT: 1
```

```
Host: localhost:42069
```

```
Connection: Keep-Alive
```

HTTP Response 2

HTTP/1.1 404 Not Found

Content-Length: 0

HTTP Request 3, a successful request of a JPG image

elton@TUESDAYSPECIAL:/mnt/c/Users/elton/Desktop/vm-shared/cs118/project1_204457607\$./webserver 42069

GET /test.jpg HTTP/1.1

Accept: text/html, application/xhtml+xml, image/jxr, */*

Accept-Language: en-US,en;q=0.5

User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/51.0.2704.79 Safari/537.36 Edge/14.14393

Accept-Encoding: gzip, deflate

DNT: 1

Host: localhost:42069

Connection: Keep-Alive

HTTP Response 3

HTTP/1.1 200 OK

Content-Type: image/jpeg

Content-Length: 116000

<<JPG data follows>>